

```

import tensorflow as tf

from tensorflow.keras import layers, datasets, models, callbacks

import matplotlib.pyplot as plt

def load_and_preprocess_data():
    """Load and preprocess MNIST data with normalization."""
    (train_images, train_labels), (test_images, test_labels) = datasets.mnist.load_data()

    train_images = train_images.reshape((train_images.shape[0], 28, 28, 1)).astype('float32') / 255
    test_images = test_images.reshape((test_images.shape[0], 28, 28, 1)).astype('float32') / 255
    return train_images, train_labels, test_images, test_labels

def build_model():
    """Build a CNN model for MNIST classification with batch normalization."""
    model = models.Sequential([
        layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
        layers.BatchNormalization(),
        layers.MaxPooling2D((2, 2)),
        layers.Conv2D(64, (3, 3), activation='relu'),
        layers.BatchNormalization(),
        layers.MaxPooling2D((2, 2)),
        layers.Conv2D(64, (3, 3), activation='relu'),
        layers.BatchNormalization(),
        layers.Flatten(),
        layers.Dense(64, activation='relu'),
        layers.Dropout(0.5), # Add dropout to prevent overfitting
        layers.Dense(10, activation='softmax')
    ])
    model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
    return model

def plot_training_history(history):

```

```
"""Plot accuracy and loss graphs for training history."""
```

```
# Plot accuracy
```

```
plt.figure(figsize=(12, 4))
```

```
plt.subplot(1, 2, 1)
```

```
plt.plot(history.history['accuracy'], label='Train')
```

```
plt.plot(history.history['val_accuracy'], label='Test')
```

```
plt.title(f'Model accuracy (final test accuracy: {history.history["val_accuracy"][-1]:.2f})')
```

```
plt.xlabel('Epochs')
```

```
plt.ylabel('Accuracy')
```

```
plt.legend(loc='upper left')
```

```
# Plot loss
```

```
plt.subplot(1, 2, 2)
```

```
plt.plot(history.history['loss'], label='Train')
```

```
plt.plot(history.history['val_loss'], label='Test')
```

```
plt.title(f'Model loss (final test loss: {history.history["val_loss"][-1]:.2f})')
```

```
plt.xlabel('Epochs')
```

```
plt.ylabel('Loss')
```

```
plt.legend(loc='upper left')
```

```
plt.tight_layout()
```

```
plt.show()
```

```
def main():
```

```
# Load and preprocess data
```

```
train_images, train_labels, test_images, test_labels = load_and_preprocess_data()
```

```
# Data augmentation
```

```
data_gen = tf.keras.preprocessing.image.ImageDataGenerator(
```

```
    rotation_range=10,
```

```
    zoom_range=0.1,
```

```
        width_shift_range=0.1,
        height_shift_range=0.1
    )

# Build model
model = build_model()

# Define callbacks
early_stopping = callbacks.EarlyStopping(monitor='val_loss', patience=3,
restore_best_weights=True)

reduce_lr = callbacks.ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=2,
min_lr=0.00001)

# Train model
history = model.fit(
    data_gen.flow(train_images, train_labels, batch_size=64),
    epochs=20,
    validation_data=(test_images, test_labels),
    callbacks=[early_stopping, reduce_lr]
)

# Evaluate model
test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f"Test accuracy: {test_acc}")

# Plot training history
plot_training_history(history)

# Save model weights
model.save('mnist_cnn_model.h5')
print("Model saved as 'mnist_cnn_model.h5'")
```

```
if __name__ == "__main__":  
    main()
```

- **Data Augmentation:**

- Defined a `data_gen` generator using `ImageDataGenerator` with parameters to rotate, zoom, and shift images. This improves generalization by adding variation to the training images.

- **Early Stopping:**

- Defined an `EarlyStopping` callback with `patience=3` to stop training if the validation loss does not improve for 3 consecutive epochs. This helps prevent overfitting.

- **Reduce Learning Rate on Plateau:**

- Used `ReduceLROnPlateau` callback to halve the learning rate if the validation loss plateaus for 2 epochs. This helps the model converge to a better minimum.

- **Batch Normalization:**


- Added `BatchNormalization` layers after each convolution layer to improve training stability and convergence speed.

Output:


```
Epoch 1/5  
/usr/local/lib/python3.10/dist-  
packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do  
not pass an `input_shape`/`input_dim` argument to a layer. When using  
Sequential models, prefer using an `Input(shape)` object as the first layer  
in the model instead.  
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)  
938/938 ————— 102s 105ms/step - accuracy:  
0.7656 - loss: 0.7411 - val_accuracy: 0.9771 - val_loss: 0.0761 -  
learning_rate: 0.0010  
Epoch 2/5  
938/938 ————— 94s 100ms/step - accuracy:  
0.9574 - loss: 0.1515 - val_accuracy: 0.9851 - val_loss: 0.0456 -  
learning_rate: 0.0010  
Epoch 3/5  
938/938 ————— 99s 105ms/step - accuracy:  
0.9687 - loss: 0.1134 - val_accuracy: 0.9798 - val_loss: 0.0683 -  
learning_rate: 0.0010  
Epoch 4/5
```

```
938/938 ————— 95s 101ms/step - accuracy:
0.9752 - loss: 0.0923 - val_accuracy: 0.9864 - val_loss: 0.0447 -
learning_rate: 0.0010
Epoch 5/5
938/938 ————— 142s 101ms/step - accuracy:
0.9766 - loss: 0.0833 - val_accuracy: 0.9919 - val_loss: 0.0275 -
learning_rate: 0.0010
313/313 ————— 4s 13ms/step - accuracy: 0.9909
- loss: 0.0340
Test accuracy: 0.9919000267982483
```

Model accuracy (final test accuracy: 0.99)



Model loss (final test loss: 0.03)



```

import tensorflow as tf

from tensorflow.keras import layers, models, optimizers

from sklearn.model_selection import train_test_split

import numpy as np

import matplotlib.pyplot as plt


def create_data(seed=42):
    """Create standardized random data for model training."""
    np.random.seed(seed)

    X = np.random.randn(1000, 10)

    y = np.random.randn(1000, 1)

    return X, y


def create_model():
    """Build a simple feedforward neural network model with an Input layer."""
    model = models.Sequential([
        layers.Input(shape=(10,)),
        layers.Dense(50, activation="relu"),
        layers.Dense(20, activation="relu"),
        layers.Dense(1)
    ])

    return model


def train_model_with_history(model, optimizer, X_train, y_train, X_val, y_val, batch_size, epochs,
optimizer_name, learning_rates):
    """Train the model with dynamically changing learning rates and track MSE and MAE for both train
and validation."""

    # Initialize dictionary to store history

    history = {
        "loss": [],
        "val_loss": [],

```

```
"mae": [],  
"val_mae": []  
}
```

```
# Loop over epochs with dynamic learning rate updates
```

```
for epoch in range(epochs):
```

```
    # Update learning rate dynamically, ensuring it is a float
```

```
    lr = float(learning_rates[epoch % len(learning_rates)])
```

```
    optimizer.learning_rate.assign(lr)
```

```
# Compile model with the updated optimizer
```

```
model.compile(optimizer=optimizer, loss="mse", metrics=["mae"])
```

```
# Train the model for one epoch and capture history
```

```
hist = model.fit(X_train, y_train, batch_size=batch_size, epochs=1, verbose=0,  
validation_data=(X_val, y_val))
```

```
# Append epoch metrics to history dictionary
```

```
history["loss"].append(hist.history["loss"][0])
```

```
history["val_loss"].append(hist.history["val_loss"][0])
```

```
history["mae"].append(hist.history["mae"][0])
```

```
history["val_mae"].append(hist.history["val_mae"][0])
```

```
# Print learning rate and losses for the epoch
```

```
print(f"Epoch {epoch+1}/{epochs} - {optimizer_name} | LR:  
{optimizer.learning_rate.numpy():.6f} | Loss: {history['loss'][-1]:.4f} | Val Loss: {history['val_loss'][-1]:.4f}")
```

```
return history
```

```
def plot_metric_history(history, metric="loss", metric_label="MSE"):
```

```
    """Plot training and validation metric history."""
```

```

plt.plot(history[metric], label=f'Train {metric_label}', color='blue')

plt.plot(history[f'val_{metric}'], label=f'Validation {metric_label}', linestyle='dashed',
color='orange')


plt.xlabel("Epochs")
plt.ylabel(metric_label)
plt.title(f'Train vs Validation {metric_label}')
plt.legend()
plt.grid(True)
plt.show()


# Create data and split into training and validation sets
X, y = create_data()
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=42)


# Initialize model
model = create_model()


# Define optimizer with an initial learning rate
optimizer = optimizers.Adam(learning_rate=0.01)


# Define learning rates to cycle through
learning_rates = [0.01, 0.005, 0.001]


# Define training parameters
epochs = 30
batch_size = 32


# Train the model
print("\nTraining with Dynamic Learning Rates:")

history = train_model_with_history(model, optimizer, X_train, y_train, X_val, y_val, batch_size,
epochs, "Adam", learning_rates)

```



```
# Plot MSE and MAE for both train and validation
```

```
plot_metric_history(history, metric="loss", metric_label="MSE")
```

```
plot_metric_history(history, metric="mae", metric_label="MAE")
```

- **Variable Learning Rates:** We define a list `learning_rates` with different rates. Each epoch cycles through these rates by updating the optimizer's learning rate.
- **Track Both Training and Validation Loss:** We split the data into training and validation sets to track `loss` and `val_loss` to monitor both overfitting and generalization.
- **Add MAE as an Additional Metric:** MAE is tracked as an additional metric by adding it to the model's metrics list during compilation, and both `mae` and `val_mae` values are stored in `history` for plotting.

Output:

```
Training with Dynamic Learning Rates:
Epoch 1/30 - Adam | LR: 0.010000 | Loss: 0.9969 | Val Loss: 1.0054
Epoch 2/30 - Adam | LR: 0.005000 | Loss: 0.9051 | Val Loss: 0.9803
Epoch 3/30 - Adam | LR: 0.001000 | Loss: 0.8715 | Val Loss: 0.9827
Epoch 4/30 - Adam | LR: 0.010000 | Loss: 0.8847 | Val Loss: 1.0452
Epoch 5/30 - Adam | LR: 0.005000 | Loss: 0.8396 | Val Loss: 1.0507
Epoch 6/30 - Adam | LR: 0.001000 | Loss: 0.7915 | Val Loss: 1.0566
Epoch 7/30 - Adam | LR: 0.010000 | Loss: 0.8349 | Val Loss: 1.0849
Epoch 8/30 - Adam | LR: 0.005000 | Loss: 0.7637 | Val Loss: 1.1152
Epoch 9/30 - Adam | LR: 0.001000 | Loss: 0.7037 | Val Loss: 1.1048
Epoch 10/30 - Adam | LR: 0.010000 | Loss: 0.7674 | Val Loss: 1.1969
Epoch 11/30 - Adam | LR: 0.005000 | Loss: 0.7203 | Val Loss: 1.1504
Epoch 12/30 - Adam | LR: 0.001000 | Loss: 0.6347 | Val Loss: 1.1591
Epoch 13/30 - Adam | LR: 0.010000 | Loss: 0.6944 | Val Loss: 1.2639
Epoch 14/30 - Adam | LR: 0.005000 | Loss: 0.6551 | Val Loss: 1.2525
Epoch 15/30 - Adam | LR: 0.001000 | Loss: 0.5730 | Val Loss: 1.2479
Epoch 16/30 - Adam | LR: 0.010000 | Loss: 0.6244 | Val Loss: 1.3787
Epoch 17/30 - Adam | LR: 0.005000 | Loss: 0.6049 | Val Loss: 1.3416
Epoch 18/30 - Adam | LR: 0.001000 | Loss: 0.4927 | Val Loss: 1.3605
Epoch 19/30 - Adam | LR: 0.010000 | Loss: 0.6048 | Val Loss: 1.4661
Epoch 20/30 - Adam | LR: 0.005000 | Loss: 0.5539 | Val Loss: 1.3734
Epoch 21/30 - Adam | LR: 0.001000 | Loss: 0.4687 | Val Loss: 1.3688
Epoch 22/30 - Adam | LR: 0.010000 | Loss: 0.5460 | Val Loss: 1.4739
Epoch 23/30 - Adam | LR: 0.005000 | Loss: 0.4758 | Val Loss: 1.4444
Epoch 24/30 - Adam | LR: 0.001000 | Loss: 0.4064 | Val Loss: 1.4443
Epoch 25/30 - Adam | LR: 0.010000 | Loss: 0.4988 | Val Loss: 1.5380
Epoch 26/30 - Adam | LR: 0.005000 | Loss: 0.4653 | Val Loss: 1.5339
Epoch 27/30 - Adam | LR: 0.001000 | Loss: 0.3762 | Val Loss: 1.4840
Epoch 28/30 - Adam | LR: 0.010000 | Loss: 0.4651 | Val Loss: 1.5686
Epoch 29/30 - Adam | LR: 0.005000 | Loss: 0.4089 | Val Loss: 1.6999
Epoch 30/30 - Adam | LR: 0.001000 | Loss: 0.3349 | Val Loss: 1.6090
```

Train vs Validation MSE

